

MARCELO VINICIUS TAPARELLI

Software Engineer | Applied AI Engineering

Campinas, SP | +55 19 99659-2477 | contato@marcelotaparelli.com.br
marcelotaparelli.com.br | LinkedIn | GitHub

PROFESSIONAL SUMMARY

Software Engineer with experience across backend, APIs, web applications, automation, and production systems. I build reliable, testable, secure software with attention to observability and performance. I deepen my Applied AI Engineering through evaluated systems and explicit controls.

PROFESSIONAL EXPERIENCE

Agência Catus | Full Stack Developer | Jul 2026 - Present

- Develop and evolve the agency's new website and two WordPress blogs, including Custom Post Types, theme customization, and new features.
- Maintain an e-commerce implementation on the Wake platform, troubleshoot production issues, and work on performance/Core Web Vitals.
- Deploy directly to server environments and continuously improve web applications.

EVAG Comunicação e Tecnologia | Developer | Apr 2024 - Present

- Developed Google Drive-to-WordPress automation for generating petition pages, reducing manual steps and making the publishing workflow more direct.
- Developed an internal request management solution integrated with GitLab, supporting organization, traceability, and flow of operational demands.
- Develop and maintain systems, websites, and automation with Laravel, PHP, Python, JavaScript, WordPress, and Tailwind, including production fixes and new features.

EVAG Comunicação e Tecnologia | Product Owner | Apr 2023 - Mar 2024

- Managed website delivery from discovery through delivery, connecting clients, design, and development.
- Structured requirements and support workflows, prioritizing across business needs, user experience, and the team's technical capacity.

TECHNICAL SKILLS

Languages: TypeScript, JavaScript, Python, PHP

Backend & runtime: Bun, Node.js, Laravel, REST APIs

Engineering: Clean Architecture, DDD principles, SOLID, Clean Code, TDD, microservices fundamentals, unit, integration and E2E tests, CI/CD, quality gates, resource limits, latency measurement, and production troubleshooting

Data & infrastructure: SQL/PostgreSQL, Redis, Docker, Terraform; AWS hands-on in a validated temporary, non-production lab: VPC/subnets, ECS/Fargate, ALB/ACM, RDS, ElastiCache, ECR, Secrets Manager, IAM, and CloudWatch; migrations, troubleshooting, validation, destroy, and audit

Security & reliability: authentication, ownership-based authorization, per-user isolation, validation, secret management, redacted logs, observability, resource limits, and regression testing

Applied AI (current focus): LLM integration, structured outputs, deterministic baselines, evals, and frozen held-out benchmarks; typed probabilistic models, exploratory confidence/calibration analysis, selective automation, hybrid policies, human review, and fallbacks; AI coding agents with human review, Python, PyTorch, and ML/DL fundamentals

Public project: Ops Triage AI — Deterministic baseline, Ollama, HybridPolicy, human review, and audit trail. Jev 1.13 was evaluated separately (evaluation-only) on the same frozen 70-ticket synthetic held-out set: 100% category accuracy, 94.29% exact tuple, 569 ms average, and US\$ 0.0038 total cost; this does not establish production superiority.

Public project: Resilient Transaction API — Bun/TS, concurrent idempotency, PostgreSQL as source of truth, Redis, retries, and circuit breaker; Terraform and temporary AWS validation, including migrations, deployment, tests, and post-destroy audit.

ADDITIONAL TRAINING

Alura: Software Engineering, Back-end, Node.js APIs, Authentication/Testing/Security, DevOps/CI/CD, Cloud/AWS, and Secure Development.

EDUCATION & LANGUAGES

Education: Systems Analysis and Development - UniBF | Completed in 2026

Languages: Portuguese (native) | English (advanced)